

# Deep Learning & Applied AI

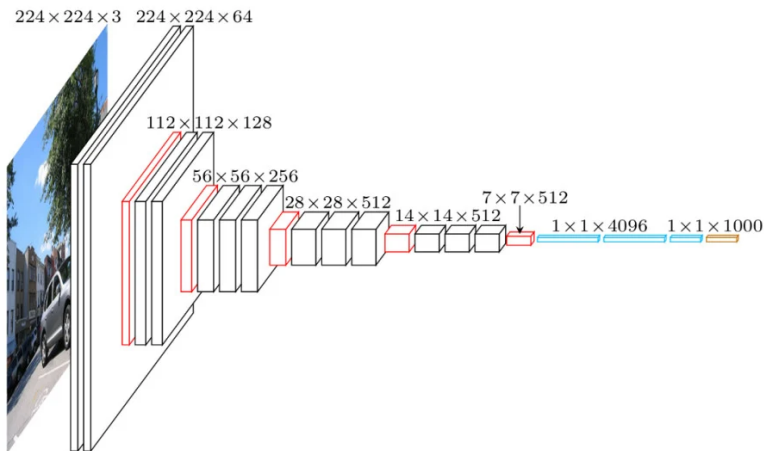
Multi-layer perceptron and back-propagation

Emanuele Rodolà  
[rodola@di.uniroma1.it](mailto:rodola@di.uniroma1.it)



# A glimpse into neural networks

In deep learning, we deal with highly parametrized models called deep neural networks:



# Deep composition

The simplest example of a nonlinear parametric model:

$$f \circ f(\mathbf{x})$$

# Deep composition

The simplest example of a nonlinear parametric model:

$$\underbrace{f}_{\text{linear}} \circ \underbrace{f}_{\text{linear}}(\mathbf{x})$$

# Deep composition

The simplest example of a nonlinear parametric model:

$$\underbrace{f \circ f}_{\text{linear}}(\mathbf{x})$$

# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

If  $\sigma$  is the logistic function, we have the [logistic regression](#) model.

# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

If  $\sigma$  is the logistic function, we have the **logistic regression** model.

Consider multiple **layers** of logistic regression models:

$$\underbrace{(\sigma \circ f)}_{\text{layer}} \circ (\sigma \circ f) \circ \cdots \circ \underbrace{(\sigma \circ f)}_{\text{layer}}(\mathbf{x})$$

# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

If  $\sigma$  is the logistic function, we have the **logistic regression** model.

Consider multiple **layers** of logistic regression models:

$$\text{output} \leftarrow \underbrace{(\sigma \circ f)}_{\text{layer } n} \circ (\sigma \circ f) \circ \cdots \circ \underbrace{(\sigma \circ f)}_{\text{layer } 1}(\mathbf{x}) \leftarrow \text{input}$$



# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

If  $\sigma$  is the logistic function, we have the **logistic regression** model.

Consider multiple **layers** of logistic regression models:

$$\text{output} \leftarrow \underbrace{(\sigma \circ f)}_{\text{output layer}} \circ (\sigma \circ f) \circ \cdots \circ \underbrace{(\sigma \circ f)}_{\text{input layer}} (\mathbf{x}) \leftarrow \text{input}$$

# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

If  $\sigma$  is the logistic function, we have the **logistic regression** model.

Consider multiple **layers** of logistic regression models:

$$\text{output} \leftarrow \underbrace{(\sigma \circ f)}_{\text{output layer}} \circ (\sigma \circ f) \circ \cdots \circ \underbrace{(\sigma \circ f)}_{\text{input layer}} (\mathbf{x}) \leftarrow \text{input}$$

(Function composition is associative, so parentheses are not necessary.)

# Deep composition

The simplest example of a nonlinear parametric model:

$$\sigma \circ f(\mathbf{x})$$

quindi una linear map seguita da una  
funzione non lineare (**activation  
function**)

If  $\sigma$  is the logistic function, we have the **logistic regression** model.

Consider multiple **layers** of logistic regression models:

$$\text{output} \leftarrow \underbrace{(\sigma \circ f)}_{\text{output layer}} \circ (\sigma \circ f) \circ \cdots \circ \underbrace{(\sigma \circ f)}_{\text{input layer}} (\mathbf{x}) \leftarrow \text{input}$$

(Function composition is associative, so parentheses are not necessary.)

More in general, consider other **activation functions** than logistic:

$$\sigma(x) = \frac{1}{1 + e^{-x}} \qquad \sigma(x) = \max\{0, x\} \quad \text{ReLU}$$

continuous

discontinuous  
gradient

# Multi-layer perceptron

We call the composition with linear  $f$  and nonlinear  $\sigma$ :

$$(\sigma \circ f) \circ (\sigma \circ f) \circ \cdots \circ (\sigma \circ f)(\mathbf{x})$$

a multi-layer perceptron (MLP) or deep feed-forward neural network.

# Multi-layer perceptron

We call the composition with linear  $f$  and nonlinear  $\sigma$ :

$$g_{\Theta}(\mathbf{x}) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x})$$

a multi-layer perceptron (MLP) or deep feed-forward neural network.

The parameters or weights of the MLP are scattered across the layers.

# Multi-layer perceptron

We call the composition with linear  $f$  and nonlinear  $\sigma$ :

$$g_{\Theta}(\mathbf{x}) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x})$$

a **multi-layer perceptron** (MLP) or **deep feed-forward neural network**.

The parameters or **weights** of the MLP are scattered across the layers.

Each layer outputs an intermediate **hidden representation**:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell})$$

where we encode the weights at layer  $\ell$  in the matrix  $\mathbf{W}_{\ell}$

# Multi-layer perceptron

We call the composition with linear  $f$  and nonlinear  $\sigma$ :

$$g_{\Theta}(\mathbf{x}) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x})$$

a **multi-layer perceptron** (MLP) or **deep feed-forward neural network**.

The parameters or **weights** of the MLP are scattered across the layers.

Each layer outputs an intermediate **hidden representation**:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell} + \mathbf{b}_{\ell})$$

where we encode the weights at layer  $\ell$  in the matrix  $\mathbf{W}_{\ell}$  and bias  $\mathbf{b}_{\ell}$ .


# Multi-layer perceptron

We call the composition with linear  $f$  and nonlinear  $\sigma$ :


$$g_{\Theta}(\mathbf{x}) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x})$$

i numeri indicano i livelli della rete

a **multi-layer perceptron** (MLP) or **deep feed-forward neural network**.

The parameters or **weights**  the MLP are scattered across the layers.

Each layer outputs an intermediate **hidden representation**:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell} + \mathbf{b}_{\ell})$$


where we encode the weights at layer  $\ell$  in the matrix  $\mathbf{W}_{\ell}$  and bias  $\mathbf{b}_{\ell}$ .

**Remark:** The **bias** can be included in the weight matrix by writing:

$$\mathbf{W} \mapsto (\mathbf{W} \quad \mathbf{b}), \quad \mathbf{x} \mapsto \begin{pmatrix} \mathbf{x} \\ 1 \end{pmatrix},$$

viene aggiunta una nuova dimensione (parametro)

because each  $f$  is **linear in the parameters** just like in linear regression.



# Hidden units

At each **hidden layer** we have:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell})$$

# Hidden units

At each **hidden layer** we have:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell})$$

Each row of the weight matrix is called a **neuron** or **hidden unit**:

$$\mathbf{W}\mathbf{x} = \begin{pmatrix} \text{— unit —} \\ \vdots \\ \text{— unit —} \end{pmatrix} \begin{pmatrix} | \\ \mathbf{x} \\ | \end{pmatrix}$$

# Hidden units

At each **hidden layer** we have:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell})$$

Each row of the weight matrix is called a **neuron** or **hidden unit**:

$$\mathbf{W}\mathbf{x} = \begin{pmatrix} \text{--- unit ---} \\ \vdots \\ \text{--- unit ---} \end{pmatrix} \begin{pmatrix} | \\ \mathbf{x} \\ | \end{pmatrix}$$

We have two interpretations:

- ① Each layer is a vector-to-vector function  $\mathbb{R}^p \rightarrow \mathbb{R}^q$ .

# Hidden units


At each **hidden layer** we have:

$$\mathbf{x}_{\ell+1} = \sigma_{\ell}(\mathbf{W}_{\ell}\mathbf{x}_{\ell})$$

Each row of the weight matrix is called a neuron or hidden unit:

$$\mathbf{W}\mathbf{x} = \begin{pmatrix} \text{--- unit ---} \\ \vdots \\ \text{--- unit ---} \end{pmatrix} \begin{pmatrix} | \\ \mathbf{x} \\ | \end{pmatrix}$$

We have two interpretations:

- 1 Each layer is a vector-to-vector function  $\mathbb{R}^p \rightarrow \mathbb{R}^q$ . 
- 2 Each layer has  $q$  units acting **in parallel**. **perché le unità sono indipendenti tra di loro**  
Each unit acts as a scalar function  $\mathbb{R}^p \rightarrow \mathbb{R}$ .


## Single layer illustration

$$\sigma(\mathbf{W}\mathbf{x}) = \sigma \circ \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \cdots & \ddots & \vdots \\ w_{m1} & w_{m2} & \cdots & w_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \sigma \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$

# Single layer illustration


$$\sigma(\mathbf{W}\mathbf{x}) = \sigma \circ \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \cdots & \ddots & \vdots \\ w_{m1} & w_{m2} & \cdots & w_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \sigma \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$

questo è un modo  
di scrivere le  
moltiplicazioni tra  
vettori di sopra  
(spoiler: è la struttura  
del MLP)




## Single layer illustration

$$\sigma(\mathbf{W}\mathbf{x}) = \sigma \circ \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \cdots & \ddots & \vdots \\ w_{m1} & w_{m2} & \cdots & w_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \sigma \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$



## Single layer illustration


$$\sigma(\mathbf{W}\mathbf{x}) = \sigma \circ \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \cdots & \ddots & \vdots \\ w_{m1} & w_{m2} & \cdots & w_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \sigma \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$





# Single layer illustration

$$\sigma(\mathbf{W}\mathbf{x}) = \sigma \circ \begin{pmatrix} w_{11} & w_{12} & \cdots & w_{1n} \\ w_{21} & w_{22} & \cdots & w_{2n} \\ \vdots & \cdots & \ddots & \vdots \\ w_{m1} & w_{m2} & \cdots & w_{mn} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} = \sigma \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix}$$



# Output layer

The output layer determines the co-domain of the network:

$$\mathbf{y} = (\sigma \circ f) \circ (\sigma \circ f) \circ \cdots \circ (\sigma \circ f)(\mathbf{x})$$

# Output layer

The output layer determines the co-domain of the network:

$$\mathbf{y} = (\sigma \circ f) \circ (\sigma \circ f) \circ \cdots \circ (\sigma \circ f)(\mathbf{x})$$

If  $\sigma$  is the logistic sigmoid, then the entire network will map:

$$\mathbb{R}^p \rightarrow (0, 1)^q$$

# Output layer

The output layer determines the co-domain of the network:

$$\mathbf{y} = (\sigma \circ f) \circ (\sigma \circ f) \circ \dots \circ (\sigma \circ f)(\mathbf{x})$$



If  $\sigma$  is the logistic sigmoid, then the entire network will map:

$$\mathbb{R}^p \rightarrow (0, 1)^q$$

→ indica un vettore di dimensione  $q$   
dove ogni dimensione può  
assumere valore  $(0, 1)$  con 0 e 1  
esclusi (classification)

For generality, it is common to have a **linear** layer at the output:

$$\mathbf{y} = f \circ (\sigma \circ f) \circ \dots \circ (\sigma \circ f)(\mathbf{x})$$

tranne nel caso di  
classificazione binaria

mapping:

$$\mathbb{R}^p \rightarrow \mathbb{R}^q$$

# Deep ReLU networks

Adding a linear layer at the output:

$$\mathbf{y} = f \circ (\sigma \circ f) \circ \cdots \circ (\sigma \circ f)(\mathbf{x})$$

# Deep ReLU networks

Adding a linear layer at the output:

$$\mathbf{y} = f \circ \sigma(\cdots)(\mathbf{x})$$

expresses  $\mathbf{y}$  as a combination of “ridge functions”  $\sigma(\cdots)$ .

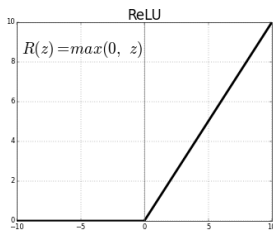
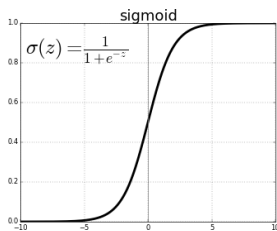
# Deep ReLU networks

Adding a linear layer at the output:

$$\mathbf{y} = f \circ \sigma(\cdots)(\mathbf{x})$$

expresses  $\mathbf{y}$  as a combination of “ridge functions”  $\sigma(\cdots)$ .

For a 2-layer network with activation  $\sigma(x) = \max\{0, x\}$  (rectifier), we get a **piecewise-linear** function:



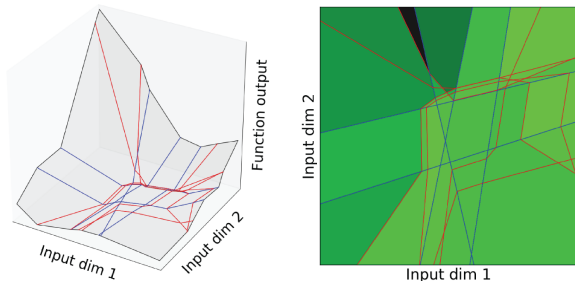
# Deep ReLU networks

Adding a linear layer at the output:

$$\mathbf{y} = f \circ \sigma(\cdots)(\mathbf{x})$$

expresses  $\mathbf{y}$  as a combination of “ridge functions”  $\sigma(\cdots)$ .

For a 2-layer network with activation  $\sigma(x) = \max\{0, x\}$  (**rectifier**), we get a **piecewise-linear** function:



The blue and red edges are produced by the **first** and **second** layer.



# Universality

What class of functions can we represent with a MLP?

# Universality

What class of functions can we represent with a MLP?

If  $\sigma$  is sigmoidal, we have the following:

**Universal Approximation Theorem** For any compact set  $\Omega \subset \mathbb{R}^p$ , the space spanned by the functions  $\phi(\mathbf{x}) = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$  is dense in  $\mathcal{C}(\Omega)$  for the uniform convergence.

# Universality

What class of functions can we represent with a MLP?

If  $\sigma$  is sigmoidal, we have the following:

**Universal Approximation Theorem** For any compact set  $\Omega \subset \mathbb{R}^p$ , the space spanned by the functions  $\phi(\mathbf{x}) = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$  is dense in  $\mathcal{C}(\Omega)$  for the uniform convergence. Thus, for any continuous function  $f$  and any  $\epsilon > 0$ , there exists  $q \in \mathbb{N}$  and weights s.t.:

$$|f(\mathbf{x}) - \sum_{k=1}^q u_k \phi(\mathbf{x})| \leq \epsilon \quad \text{for all } \mathbf{x} \in \Omega$$

# Universality

What class of functions can we represent with a MLP?

If  $\sigma$  is sigmoidal, we have the following:

**Universal Approximation Theorem** For any compact set  $\Omega \subset \mathbb{R}^p$ , the space spanned by the functions  $\phi(\mathbf{x}) = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$  is dense in  $\mathcal{C}(\Omega)$  for the uniform convergence. Thus, for any continuous function  $f$  and any  $\epsilon > 0$ , there exists  $q \in \mathbb{N}$  and weights s.t.:

$$|f(\mathbf{x}) - \sum_{k=1}^q u_k \phi(\mathbf{x})| \leq \epsilon \quad \text{for all } \mathbf{x} \in \Omega$$

The network in the theorem has just **one** hidden layer.

# Universality

What class of functions can we represent with a MLP?

If  $\sigma$  is sigmoidal, we have the following:

**Universal Approximation Theorem** For any compact set  $\Omega \subset \mathbb{R}^p$ , the space spanned by the functions  $\phi(\mathbf{x}) = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$  is dense in  $\mathcal{C}(\Omega)$  for the uniform convergence. Thus, for any continuous function  $f$  and any  $\epsilon > 0$ , there exists  $q \in \mathbb{N}$  and weights s.t.:

$$|f(\mathbf{x}) - \sum_{k=1}^q u_k \phi(\mathbf{x})| \leq \epsilon \quad \text{for all } \mathbf{x} \in \Omega$$

The network in the theorem has just **one** hidden layer.

For large enough  $q$ , the training error can be made **arbitrarily small**.

# Training

Given a MLP with training pairs  $\{\mathbf{x}_i, \mathbf{y}_i\}$ :

$$g_{\Theta}(\mathbf{x}_i) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x}_i) = \mathbf{y}_i$$

# Training

Given a MLP with training pairs  $\{\mathbf{x}_i, \mathbf{y}_i\}$ :

$$g_{\Theta}(\mathbf{x}_i) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x}_i) = \mathbf{y}_i$$

Consider the MSE loss:

$$\ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{y}_i - g_{\Theta}(\mathbf{x}_i)\|_2^2$$

Solving for the weights  $\Theta$  is referred to as **training**.

# Training

Given a MLP with training pairs  $\{\mathbf{x}_i, \mathbf{y}_i\}$ :

$$g_{\Theta}(\mathbf{x}_i) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x}_i) = \mathbf{y}_i$$

Consider the MSE loss:

$$\ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{y}_i - g_{\Theta}(\mathbf{x}_i)\|_2^2$$

Solving for the weights  $\Theta$  is referred to as **training**.

In general, the loss is **not convex** w.r.t.  $\Theta$ .



# Training

Given a MLP with training pairs  $\{\mathbf{x}_i, \mathbf{y}_i\}$ :

$$g_{\Theta}(\mathbf{x}_i) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \cdots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x}_i) = \mathbf{y}_i$$

Consider the MSE loss:

$$\ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{y}_i - g_{\Theta}(\mathbf{x}_i)\|_2^2$$

Solving for the weights  $\Theta$  is referred to as **training**.

In general, the loss is **not convex** w.r.t.  $\Theta$ .

As we have seen, the following **special cases** are convex:

- One layer, no activation, MSE loss ( $\Rightarrow$  linear regression).

# Training

Given a MLP with training pairs  $\{\mathbf{x}_i, \mathbf{y}_i\}$ :

$$g_{\Theta}(\mathbf{x}_i) = (\sigma \circ f_{\Theta_n}) \circ (\sigma \circ f_{\Theta_{n-1}}) \circ \dots \circ (\sigma \circ f_{\Theta_1})(\mathbf{x}_i) = \mathbf{y}_i$$

Consider the MSE loss:

$$\ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \|\mathbf{y}_i - g_{\Theta}(\mathbf{x}_i)\|_2^2$$

Solving for the weights  $\Theta$  is referred to as **training**.

**In general, the loss is not convex w.r.t.  $\Theta$ .**

As we have seen, the following **special cases** are convex:

- One layer, no activation, MSE loss ( $\Rightarrow$  linear regression).
- One layer, sigmoid activation, logistic loss ( $\Rightarrow$  logistic regression).

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|\mathbf{y}_i - g_{\Theta}(\mathbf{x}_i)\|_2^2$$

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|\mathbf{y}_i - (\sigma(f_{\Theta_n}((\sigma(f_{\Theta_{n-1}}(\cdots(\sigma(f_{\Theta_1}(\mathbf{x}_i))\cdots))))))\|_2^2$$

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|\mathbf{y}_i - (\sigma(f_{\Theta_n}((\sigma(f_{\Theta_{n-1}}(\cdots(\sigma(f_{\Theta_1}(\mathbf{x}_i))\cdots))))))\|_2^2$$

- Computing the gradients **by hand** is infeasible.

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|(\mathbf{y}_i - (\sigma(f_{\Theta_n}((\sigma(f_{\Theta_{n-1}}(\cdots(\sigma(f_{\Theta_1}(\mathbf{x}_i)) \cdots))))))\|_2^2$$

- Computing the gradients **by hand** is infeasible.
- **Finite differences** require  $O(\# \text{weights})$  evaluations of  $\ell_{\Theta}$ .

# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|(\mathbf{y}_i - (\sigma(f_{\Theta_n}((\sigma(f_{\Theta_{n-1}}(\cdots(\sigma(f_{\Theta_1}(\mathbf{x}_i)) \cdots))))))\|_2^2$$

- Computing the gradients **by hand** is infeasible.
- **Finite differences** require  $O(\#\text{weights})$  evaluations of  $\ell_{\Theta}$ .
- Using the **chain rule** is sub-optimal.



# Training

We train using gradient descent-like algorithms.

**Bottleneck:** Computation of gradients  $\nabla \ell_{\Theta}$ .

For the basic MSE, this means:

$$\nabla \ell_{\Theta}(\{\mathbf{x}_i, \mathbf{y}_i\}) = \frac{1}{n} \sum_{i=1}^n \nabla_{\Theta} \|(\mathbf{y}_i - (\sigma(f_{\Theta_n}((\sigma(f_{\Theta_{n-1}}(\cdots(\sigma(f_{\Theta_1}(\mathbf{x}_i))\cdots))))))\|_2^2$$

- Computing the gradients **by hand** is infeasible.
- **Finite differences** require  $O(\#\text{weights})$  evaluations of  $\ell_{\Theta}$ .
- Using the **chain rule** is sub-optimal.

We want to automatize this **computational step** efficiently.

# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$

$x$   
●


# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$




# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$




# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



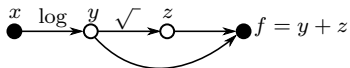
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$

$x$   
●


# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$






# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.


**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



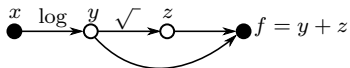
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

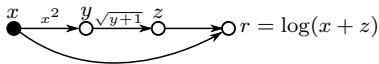
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



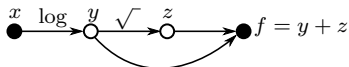
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

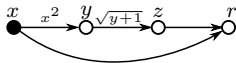
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



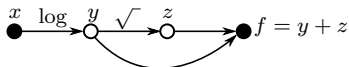
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

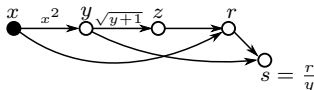
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



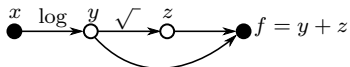
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.


**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



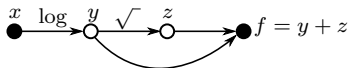
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

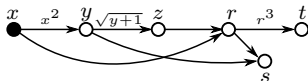
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



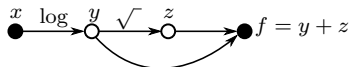
# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

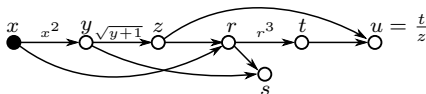
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$




# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .

A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

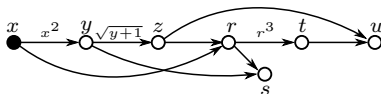
**Example:**

$$f(x) = \log x + \sqrt{\log x}$$



**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$





# Computational graphs

Consider a generic function  $f : \mathbb{R} \rightarrow \mathbb{R}$ .


A **computational graph** is a directed acyclic graph representing the computation of  $f(x)$  with **intermediate** variables.

le variabili intermedie sono i pallini bianchi, quelli neri sono l'input e l'output

**Example:**


$$f(x) = \log x + \sqrt{\log x}$$

le frecce indicano che l'operazione è salvata nella variabile successiva



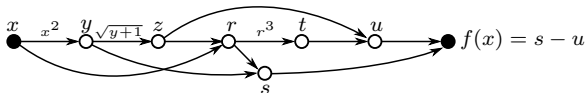
**Example:**

$$f(x) = \frac{\log(x + \sqrt{x^2 + 1})}{x^2} - \frac{\log^3(x + \sqrt{x^2 + 1})}{\sqrt{x^2 + 1}}$$



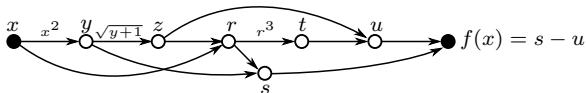
# Computational graphs

The evaluation of  $f(x)$  corresponds to a **forward traversal** of the graph:



# Computational graphs

The evaluation of  $f(x)$  corresponds to a **forward traversal** of the graph:

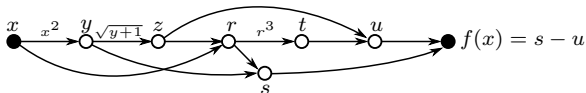


The graph is constructed programmatically, for example:

```
z = sqrt(sum(square(x), 1));
```

# Computational graphs

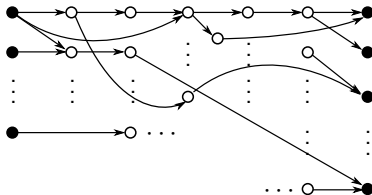
The evaluation of  $f(x)$  corresponds to a **forward traversal** of the graph:



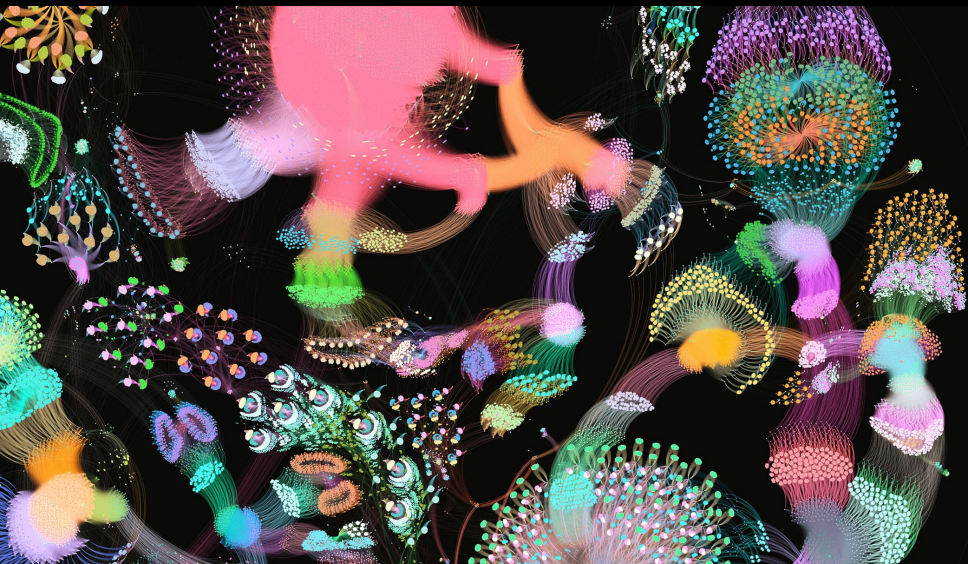
The graph is constructed programmatically, for example:

```
z = sqrt(sum(square(x), 1));
```

For **high-dimensional** input/output, the graph may be more complex:



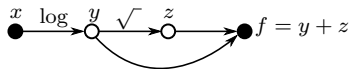
The computational graph gets big quickly.



Poplar visualization, see <https://www.graphcore.ai/products/poplar>

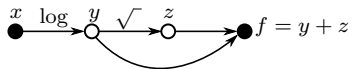
# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



# Automatic differentiation: Forward mode

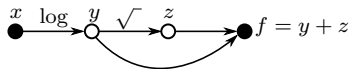
$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial x}{\partial x} = 1$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



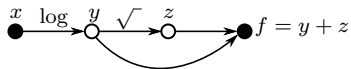
$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x}$$



# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$

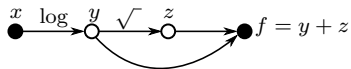


$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$




$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



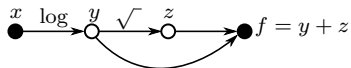
$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



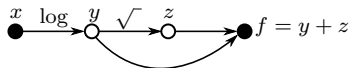
$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



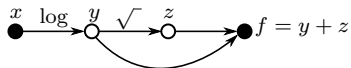
$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x} = \frac{1}{2\sqrt{y}} \frac{\partial y}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial x}{\partial x} = 1$$

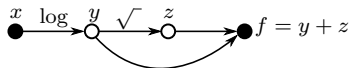
$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x} = \frac{1}{2\sqrt{y}} \frac{\partial y}{\partial x}$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial x}{\partial x} = 1$$

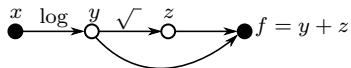
$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x} = \frac{1}{2\sqrt{y}} \frac{\partial y}{\partial x}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x} + \frac{\partial f}{\partial z} \frac{\partial z}{\partial x}$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x} = \frac{1}{2\sqrt{y}} \frac{\partial y}{\partial x}$$


$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x} + \frac{\partial f}{\partial z} \frac{\partial z}{\partial x} = \frac{\partial (y + z)}{\partial y} \frac{\partial y}{\partial x} + \frac{\partial (y + z)}{\partial z} \frac{\partial z}{\partial x}$$



# Automatic differentiation: Forward mode

## PRIMO METODO PER FARE AUTOMATIC DIFFERENTIATION

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial x}{\partial x} = 1$$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial x} \frac{\partial x}{\partial x} = \frac{\partial \log x}{\partial x} \frac{\partial x}{\partial x} = \frac{1}{x} \frac{\partial x}{\partial x}$$

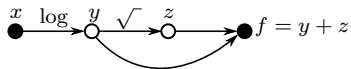
$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial \sqrt{y}}{\partial y} \frac{\partial y}{\partial x} = \frac{1}{2\sqrt{y}} \frac{\partial y}{\partial x}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x} + \frac{\partial f}{\partial z} \frac{\partial z}{\partial x} = \frac{\partial (y + z)}{\partial y} \frac{\partial y}{\partial x} + \frac{\partial (y + z)}{\partial z} \frac{\partial z}{\partial x} = \frac{\partial y}{\partial x} + \frac{\partial z}{\partial x}$$

in pratica, calcola la derivata parziale  
per ogni nodo (variabile intermedia e output)  
SULLA BASE DI X (l'input) applicando la  
chain rule e poi somma tutto. Questo è il  
significato di "forward mode"

# Automatic differentiation: Forward mode

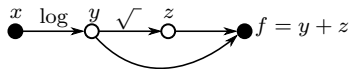
$$f(x) = \log x + \sqrt{\log x} \qquad \frac{\partial f}{\partial x} = \frac{1}{2\sqrt{y}} \frac{1}{x} + \frac{1}{x}$$



Assumption: Each partial derivative is a “primitive” accessible in **closed form** and can be computed on the fly.

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x} \qquad \frac{\partial f}{\partial x} = \frac{1}{2\sqrt{y}} \frac{1}{x} + \frac{1}{x}$$



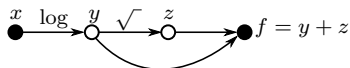
Assumption: Each partial derivative is a “primitive” accessible in **closed form** and can be computed on the fly.


$$\text{cost of computing } \frac{\partial f}{\partial x}(x) = \text{cost of computing } f(x)$$

# Automatic differentiation: Forward mode

$$f(x) = \log x + \sqrt{\log x} \qquad \frac{\partial f}{\partial x} = \frac{1}{2\sqrt{y}} \frac{1}{x} + \frac{1}{x}$$

era il risultato della slide precedente



Assumption: Each partial derivative is a “primitive” accessible in **closed form** and can be computed on the fly. 

$$\text{cost of computing } \frac{\partial f}{\partial x}(x) = \text{cost of computing } f(x)$$

However, if the input is high-dimensional, i.e.  $f : \mathbb{R}^p \rightarrow \mathbb{R}$ :

$$\text{cost of computing } \nabla f(\mathbf{x}) = p \times \text{cost of computing } f(\mathbf{x})$$

since partial derivatives must be computed w.r.t. each input dimension.

In pratica, è efficiente per funzioni con un piccolo numero di input

# Automatic differentiation: Forward mode

The forward mode computes all the partial derivatives  $\frac{\partial y}{\partial x}, \frac{\partial z}{\partial x}, \dots$  with respect to the **input**  $x$ .

Straightforward application of the **chain rule**.

# Automatic differentiation: Forward mode

The forward mode computes all the partial derivatives  $\frac{\partial y}{\partial x}, \frac{\partial z}{\partial x}, \dots$  with respect to the **input**  $x$ .

Straightforward application of the **chain rule**.

Automatic differentiation $\neq$ Symbolic differentiation
(e.g. autograd) (e.g. Mathematica)

We accumulate values during code execution to generate numerical derivative **evaluations** rather than derivative **expressions**.

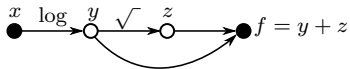
# Automatic differentiation: Forward mode

The forward mode computes all the partial derivatives  $\frac{\partial y}{\partial x}, \frac{\partial z}{\partial x}, \dots$  with respect to the **input**  $x$ .

Straightforward application of the **chain rule**.

Automatic differentiation  $\neq$  Symbolic differentiation  
(e.g. autograd) (e.g. Mathematica)

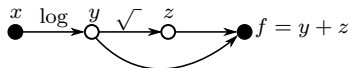
We accumulate values during code execution to generate numerical derivative **evaluations** rather than derivative **expressions**.



**Reverse mode**: compute all the partial derivatives  $\frac{\partial f}{\partial z}, \dots, \frac{\partial f}{\partial x}$  with respect to the **inner nodes**.

# Automatic differentiation: Reverse mode

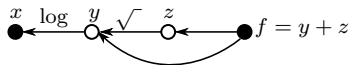
$$f(x) = \log x + \sqrt{\log x}$$





# Automatic differentiation: Reverse mode

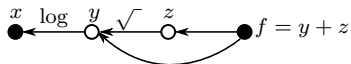
$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial x} = 1$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

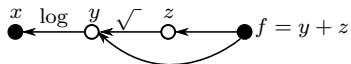
$$\frac{\partial f}{\partial z} =$$

$$\frac{\partial f}{\partial y} =$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

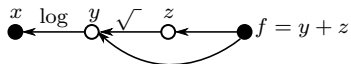
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z}$$

$$\frac{\partial f}{\partial y} =$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

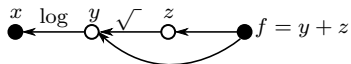
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z}$$

$$\frac{\partial f}{\partial y} =$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

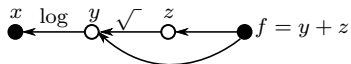
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial(y+z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} =$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

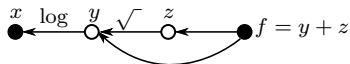
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial f}{\partial y}$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

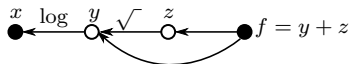
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial y}$$

$$\frac{\partial f}{\partial x} =$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$

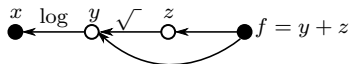
$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial y} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} =$$



# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

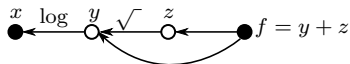
$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial y} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x}$$

# Automatic differentiation: Reverse mode

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$


$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial y} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial \log x}{\partial x}$$

# Automatic differentiation: Reverse mode

## SECONDO METODO PER FARE AUTOMATIC DIFFERENTIATION

$$f(x) = \log x + \sqrt{\log x}$$



$$\frac{\partial f}{\partial f} = 1$$

viene ripetuto lo stesso procedimento di prima, ma inversamente (sulla base dei nodi intermedi partendo da  $f$ )

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial z}{\partial y} + \frac{\partial f}{\partial y} \frac{\partial f}{\partial y} \frac{\partial \sqrt{y}}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y + z)}{\partial y} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial \log x}{\partial x} = \frac{\partial f}{\partial y} \frac{1}{x}$$

# Automatic differentiation: Reverse mode

Reverse mode requires computing the values of the **internal nodes** first:

$$\frac{\partial f}{\partial f} = 1$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial f} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial \log x}{\partial x} = \frac{\partial f}{\partial y} \frac{1}{x}$$

# Automatic differentiation: Reverse mode

Reverse mode requires computing the values of the **internal nodes** first:

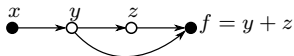
$$\frac{\partial f}{\partial f} = 1$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial f} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial \log x}{\partial x} = \frac{\partial f}{\partial y} \frac{1}{x}$$

- ① **Forward pass** to evaluate all the interior nodes  $y, z, \dots$



**Remark:** This is **not** forward-mode autodiff, since we are only computing **function values**.

# Automatic differentiation: Reverse mode


Reverse mode requires computing the values of the **internal nodes** first:

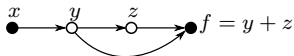
$$\frac{\partial f}{\partial f} = 1$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial z} = \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial z} \frac{\partial \sqrt{y}}{\partial y} + \frac{\partial f}{\partial f} \frac{\partial (y+z)}{\partial f} = \frac{\partial f}{\partial z} \frac{1}{2\sqrt{y}} + \frac{\partial f}{\partial f}$$

$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial \log x}{\partial x} = \frac{\partial f}{\partial y} \frac{1}{x}$$

- ① Forward pass  evaluate all the interior nodes  $y, z, \dots$



**Remark:** This is not forward-mode autodiff, since we are only computing function values.

- ② Backward pass to compute the derivatives.



La reverse mode è più efficiente per funzioni con un grande numero di parametri (come nelle reti neurali)

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

il codominio di una loss function  
sarà sempre unidimensionale

where  $p \gg 1$  is the number of **weights**.

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

matrici di gradienti

Instead of simple derivatives we must compute **gradients** and **Jacobians**.



# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

$$\ell = \epsilon(\sigma \circ f \circ \sigma \circ f \circ \dots \circ f)$$

$\epsilon$  computes the actual scalar error for the loss.

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

$$\ell = \epsilon(f_{t-1} \circ f_{t-2} \circ \cdots \circ f_2 \circ f_1)$$

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

$$\ell = \epsilon(f_{t-1} \circ f_{t-2} \circ \cdots \circ f_2 \circ f_1)$$

Denote by  $\mathbf{J}_k$  the Jacobian at layer  $k$ .

- Forward-mode autodiff:

$$\nabla \ell = \mathbf{J}_{t-1}(\mathbf{J}_{t-2}(\cdots (\mathbf{J}_3(\mathbf{J}_2 \mathbf{J}_1))))$$

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

$$\ell = \epsilon(f_{t-1} \circ f_{t-2} \circ \cdots \circ f_2 \circ f_1)$$

Denote by  $\mathbf{J}_k$  the Jacobian at layer  $k$ .

- Forward-mode autodiff:

$$\nabla \ell = \mathbf{J}_{t-1}(\mathbf{J}_{t-2}(\cdots (\mathbf{J}_3(\mathbf{J}_2 \mathbf{J}_1))))$$

- Reverse-mode autodiff:

$$\nabla \ell = (((((\mathbf{J}_{t-1} \mathbf{J}_{t-2}) \mathbf{J}_{t-3}) \cdots) \mathbf{J}_2) \mathbf{J}_1$$

# Back-propagation

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

$$\ell = \epsilon(f_{t-1} \circ f_{t-2} \circ \cdots \circ f_2 \circ f_1)$$

Denote by  $\mathbf{J}_k$  the Jacobian at layer  $k$ .

- **Forward**-mode autodiff:

$$\nabla \ell = \mathbf{J}_{t-1}(\mathbf{J}_{t-2}(\cdots(\mathbf{J}_3(\mathbf{J}_2\mathbf{J}_1)))) \quad \# \text{ ops: } p \sum_{k=2}^{t-1} d_k d_{k+1}$$

- Reverse-mode autodiff:

$$\nabla \ell = (((\mathbf{J}_{t-1}\mathbf{J}_{t-2})\mathbf{J}_{t-3})\cdots)\mathbf{J}_2)\mathbf{J}_1$$

# Back-propagation ridenominazione della reverse mode

When training neural nets, we compute the gradient of a loss

$$\ell : \mathbb{R}^p \rightarrow \mathbb{R}^1$$

where  $p \gg 1$  is the number of **weights**.

Instead of simple derivatives we must compute **gradients** and **Jacobians**.

le derivate  
allo step  $k$

$$\ell = \epsilon(f_{t-1} \circ f_{t-2} \circ \cdots \circ f_2 \circ f_1)$$

Denote by  $\mathbf{J}_k$  the Jacobian at layer  $k$ .

• **Forward**-mode autodiff:

$$\nabla \ell = \mathbf{J}_{t-1}(\mathbf{J}_{t-2}(\cdots(\mathbf{J}_3(\mathbf{J}_2\mathbf{J}_1)))) \quad \# \text{ ops: } p \sum_{k=2}^{t-1} d_k d_{k+1}$$

• **Reverse**-mode autodiff:

$$\nabla \ell = (((((\mathbf{J}_{t-1}\mathbf{J}_{t-2})\mathbf{J}_{t-3})\cdots)\mathbf{J}_2)\mathbf{J}_1) \quad \# \text{ ops: } 1 \sum_{k=1}^{t-2} d_k d_{k+1}$$

le parentesi indicano che  
il risultato di quella chain rule  
è salvato in qualche modo

# Back-propagation

We call **back-propagation** the reverse mode automatic differentiation applied to deep neural networks.

Evaluating  $\nabla \ell$  with backprop is as fast as evaluating  $\ell$ .

# Back-propagation

We call **back-propagation** the reverse mode automatic differentiation applied to deep neural networks.

Evaluating  $\nabla \ell$  with backprop is as fast as evaluating  $\ell$ .

Back-propagation is not just the chain rule.



# Back-propagation

We call **back-propagation** the reverse mode automatic differentiation applied to deep neural networks.

Evaluating  $\nabla \ell$  with backprop is as fast as evaluating  $\ell$ .

Back-propagation is not just the chain rule.

In fact, not even the costly forward mode is just the chain rule. There are **intermediate variables**. Backprop is a **computational** technique.

# Back-propagation

We call **back-propagation** the reverse mode automatic differentiation applied to deep neural networks.

Evaluating  $\nabla \ell$  with backprop is as fast as evaluating  $\ell$ .

Back-propagation is not just the chain rule.

In fact, not even the costly forward mode is just the chain rule. There are **intermediate variables**. Backprop is a **computational** technique.

Backprop through computational graph of the loss

$\approx$

Backprop “through the network”

## Suggested reading

Nice, accessible survey on automatic differentiation:

<https://arxiv.org/pdf/1502.05767>